

When Models Debate Themselves

Argument vs. arguer in LLM-on-LLM evaluation

Final-project report — multi-agent LangGraph study

Abstract

Multi-agent debate has emerged as a candidate evaluation pattern for large language models — pit two models against each other and have a third judge the resulting arguments. The same pattern underlies *AI safety via debate* (Irving et al. 2018) and a growing body of LLM-as-judge protocols. Before those uses can be trusted, a basic prerequisite must hold: the verdict has to track *argument content* rather than the identity of the model that produced it.

We built a six-node LangGraph multi-agent debate system in which two LLMs argue opposing sides of the resolution *the human brain is smarter than any large language model*. A third model (Haiku 4.5) summarizes, and a fourth (Opus 4.7) issues the verdict. We ran a 2×2 factorial — *model assignment* (Opus arguing PRO vs Sonnet arguing PRO) crossed with the *presence of a shared evidence pack* (87 facts drawn from cognitive- science and machine-learning literature, identical for both sides). Across all four runs the side argued by Claude Opus 4.7 won or tied; in no run did it lose. The evidence pack narrowed but did not eliminate the effect — it converted the Opus-CON win in the no-pack condition into a DRAW in the pack condition. We interpret the finding under two compatible registers — a *persona effect* familiar from the cognitive psychology of argumentation, and *judge bias / family preference* familiar from the LLM-as-judge literature. Implications for evaluation methodology, design choices that mitigate the effect, and follow-up experiments are discussed.

Keywords — multi-agent debate; LLM-as-judge; LangGraph; cognitive science; argumentation; symbol grounding; shared retrieval; evaluation.

1. Introduction

1.1 Motivation

The default way to evaluate a language model has, over the past two years, shifted from human ratings toward LLM-graded benchmarks. *MT-Bench* and *Chatbot Arena* (Zheng et al. 2023), *AlpacaEval* (Li et al. 2023b), and a long tail of in-house preference-ranking pipelines all rely on one model grading another. The cost reduction is dramatic: a 1 000-prompt evaluation that would have taken an expert pool a week now takes ten minutes of API time. The methodological cost is less explored.

A separate line of research — *AI safety via debate* (Irving, Christiano & Amodei 2018) — proposes the structurally similar idea that two AI systems should argue opposing sides of a question with a human or a third AI as judge. The hope is that, even when individual claims are too long or too technical for a

judge to verify directly, the dialectical pressure of two sides surfaces the truth. This is a strong claim. It depends on the judge actually evaluating the *arguments* rather than the *arguers*.

1.2 Research question

We ask, in the cleanest form we could construct:

Does the model assignment of the debaters dominate the verdict, or does argument content survive equally across model swaps?

The question splits naturally into two hypotheses.

- **H1 (null).** With identical prompts, identical topic, and identical time budgets, swapping which model argues PRO vs CON does not change the winner. *If H1 holds, the system is reading content.*
- **H2 (bridge).** Adding a shared, symmetric evidence pack to both sides' system prompts neutralizes any model effect that exists. *If H2 holds and H1 fails, the system is style-driven but evidence-correctable.*

1.3 Contributions

This report makes three contributions.

1. A small, reproducible LangGraph harness (~250 lines of Python) that instantiates a four-round, eight-minute multi-agent debate with structured logging and a shared-evidence-pack injection point. The harness is model-agnostic and isolates one experimental variable at a time.
2. A 2×2 factorial result against the resolution *the human brain is smarter than any LLM*. Whichever side Opus 4.7 argues wins or ties in four of four runs.
3. A bridge interpretation that connects the empirical finding to two adjacent literatures — the cognitive psychology of argumentation and the LLM-as-judge methodology critique — and proposes concrete follow-ups that distinguish the two readings.

2. Background

2.1 The cognitive science of debate

A debating mind draws on three subsystems.

Memory. Working memory holds about four chunks at any moment (Cowan 2001), down from Miller's (1956) classic 7 ± 2 . Baddeley's multi-component model decomposes that capacity into a phonological loop, a visuospatial sketchpad, an episodic buffer, and a central executive. Long-term memory is partitioned (Tulving 1985) into episodic, semantic, and procedural stores, each on distinct neural substrates. Encoding follows depth-of-processing (Craik & Lockhart 1972): semantic encoding produces stronger recall than surface encoding. Retrieval is reconstructive rather than playback (Loftus 1974) — humans rebuild memories at retrieval time, which is one reason eyewitness testimony is unreliable.

Reasoning. Kahneman's (2011) dual-process framing distinguishes fast intuitive cognition (System 1) from slow deliberative cognition (System 2). Causal-model construction starts in toddlerhood (Gopnik 2012): pre-linguistic children run interventionist experiments on objects. Concepts are grounded in sensorimotor experience — Harnad's (1990) symbol-grounding problem; Barsalou's (1999) perceptual

symbol systems; Lakoff & Johnson's (1980) embodied metaphor. Symbols pointing only to other symbols do not carry meaning until they touch a non-symbolic referent.

Metacognitive monitoring. Humans have explicit feeling-of-knowing, confidence judgments, and source monitoring. They are not always good at this — Dunning-Kruger (1999) describes systematic mis-calibration in low-skill individuals — but the *capacity* for self-modeling is present and can be invoked. Documented biases (Tversky & Kahneman 1974) include anchoring, framing, availability, and hindsight; fatigue effects modulate deliberation (Danziger et al. 2011, with effect-size since debated).

2.2 The computer science of LLMs

A modern frontier LLM is a transformer (Vaswani et al. 2017) trained auto-regressively on roughly 10^{13} tokens with 10^{11} – 10^{12} parameters. Self-attention lets each token attend to every other token in the context window — content-addressable retrieval at $O(n^2)$ cost. Multi-head attention parallelizes that across subspaces. Decoder-only models predict the next token; their loss is cross-entropy on next-token prediction. Everything called "reasoning" or "argumentation" emerges from scaling that single objective and adding two post-training stages — supervised fine-tuning on human demonstrations, then RLHF or DPO for preference alignment (Ouyang et al. 2022; Bai et al. 2022).

The natural mapping to human cognition is partial but instructive (§2.3). The working-memory analogue is the context window — generous in size (200 K – 2 M tokens for current frontier models) but vulnerable to the *lost-in-the-middle* effect (Liu et al. 2023): retrieval accuracy is highest at the start and end of the context, lower in the middle. The long-term-memory analogue is the parameter set itself, frozen at inference. There is no native consolidation: new facts require fine-tuning or retrieval-augmented generation (RAG; Lewis et al. 2020).

Documented limits matter for any debate study. Mirzadeh et al. (2024) introduced *GSM-Symbolic*, a math-word-problem benchmark in which renaming variables and numbers drops frontier-model accuracy by 30–65 %. Turpin et al. (2023) showed that chain-of-thought is often *unfaithful* — verbalized reasoning does not reflect the actual computation driving the answer. Kadavath et al. (2022) found that calibration is reasonable in- distribution and breaks down out-of-distribution. Sharma et al. (2023) document sycophancy emerging from RLHF.

2.3 The mapping

The two views align surprisingly well — same primitives, different substrates.

Working memory	≈ 4 chunks (Cowan); seconds without rehearsal	200 K – 2 M token context window; volatile per-session
Long-term memory	episodic / semantic / procedural (Tulving)	parametric weights from pretraining; no episodic store
Encoding	depth-of-processing; semantic > surface	distributed embeddings in $\mathbb{R}^d$
Retrieval	cue-dependent, reconstructive (Tulving, Loftus)	soft attention; RAG for external memory

Learning input	$\approx 10^7$ words by age 5	$\approx 10^{13}$ tokens during pretraining
Reasoning	System 1 / System 2 (Kahneman)	next-token prediction + chain-of-thought + RLHF
Grounding	sensorimotor (Harnad, Barsalou, Lakoff)	token co-occurrence; second-order at best
Calibration	explicit metacognition; Dunning-Kruger	in-distribution good (Kadavath); OOD breakdown
Energy budget	≈ 20 W continuous	tens of MW for training; multi-MW inference fleets

Both literatures address the same primitives with different tools, which is what makes a fair experimental comparison possible.

3. Literature review

3.1 The cognitive psychology of argumentation

Argumentation is older than psychology. Toulmin (1958) introduced the field's standard schema (claim, data, warrant, backing, rebuttal, qualifier) and that schema still informs how human judges score debates. Walton's pragma-dialectic tradition (e.g., Walton 1996) extends Toulmin into argument schemes and presumptive reasoning. The relevance for our study is direct: human debate judging is a cognitive task with a known structure, and that structure can be mimicked in prompt design.

The most provocative recent line is Mercier & Sperber's (2011, 2017) *argumentative theory of reasoning*. Their thesis: human reasoning did not evolve to track truth in solitary cognition — it evolved to *win arguments*. Reasoning is a social organ. The signature predictions — confirmation bias, motivated reasoning, the lawyer-not-scientist asymmetry — are confirmed in laboratory work. Two implications matter here. First, debate is the natural setting in which reasoning works at its best. Second, deliberators tend to *over-weight argumentation style* — citation density, calibrated hedging, structural completeness — because in social settings those features predict trustworthy interlocutors. This becomes a candidate explanation for the model effect we observe in §6.

There is also a long tradition of work on *judges* of argumentation. Petty & Cacioppo's (1986) elaboration-likelihood model distinguishes central-route (argument-content) from peripheral-route (source-cue) processing. Deliberative readers are presumed to be on the central route, but laboratory work consistently finds peripheral cues — expertise signals, hedging style, formatting — leaking into nominally-central judgments. Our LLM judge sits in this same position.

3.2 LLM-as-judge — known biases

The empirical literature on LLM judges is now substantial. Zheng et al. (2023) introduced MT-Bench and Chatbot Arena and documented two robust biases: a *position bias* in which the model preferentially favors the first answer presented, and a *length bias* in which longer answers are rated higher. Wang et al. (2023, *Large Language Models Are Not Fair Evaluators*) extended the catalog with verbosity and

self-preference biases.

Self-preference — a model rating its own output more favorably than a peer's — is the bias most relevant here. Panickssery et al. (2024) demonstrate that frontier models show statistically significant self-preference even when blinded to authorship, and that this preference strengthens for stylistic regularities the model itself produces. If the judge in our study (Opus 4.7) shares stylistic regularities with one of the debaters (also Opus 4.7), self-preference predicts an Opus-side bias. This is exactly the pattern we report. We cannot, with our four-run sample, distinguish self-preference from a more general "deliberative register" preference; we discuss this in §8 (Limitations) and §9 (Future work).

A complementary line of work studies LLM-judge *agreement* with humans. On well-defined tasks (math, coding) GPT-4-class judges agree with expert humans 80–90 % of the time; on subjective tasks (debate judgment, creative writing) agreement falls. This is consistent with the persona-effect reading: when the judgment is subjective, peripheral cues do more work.

3.3 Multi-agent debate for LLMs

Irving, Christiano & Amodei (2018) is the foundational *AI safety via debate* paper. The proposal: two AI systems argue, a (typically weaker) judge picks a winner, and the equilibrium of the game incentivizes honesty even when individual claims are too long or technical to verify. The proposal makes strong assumptions on judge competence.

Du et al. (2023, *Improving Factuality and Reasoning via Multi-Agent Debate*) showed empirically that multi-round debate among LLM agents improves accuracy on math and reasoning benchmarks, even without an external judge — agents converge on correct answers more often than single-turn or self-consistency baselines. Liang et al. (2023) extended this with *encouraging divergent thinking* in agents. Khan et al. (2024) directly tested the safety-via-debate proposal: with a human or a weaker LLM judge, *more persuasive* debaters lead to *more truthful* judgments under their setup. This is encouraging — but it presupposes that the judge is evaluating arguments rather than arguer-style, which is the question we ask here.

3.4 Retrieval augmentation and shared grounding

Lewis et al. (2020) introduced retrieval-augmented generation (RAG): augment the LLM's context with retrieved passages from an external corpus. Two well-established findings carry over to our setup. First, RAG substantially closes the gap between in-distribution and out-of-distribution performance — supplying relevant facts at inference time partially compensates for what the parameters do not encode. Second, the gap is *closed*, not *eliminated* — model capability still modulates how retrieved evidence is integrated. Our materials pack is effectively a shared, in-context retrieval source identical for both sides. The §6 finding that materials *narrow but do not eliminate* the model effect mirrors the RAG literature.

A useful negative result: Liu et al. (2023) showed that simply giving the model more context does not always help — relevant passages buried in the middle of a long context are reliably under-utilized. We position our briefing at the *start* of the system prompt and keep it under 3 000 tokens to avoid this failure mode.

3.5 The gap this study addresses


To our knowledge, the model-vs-content question has not been asked in the specific setup we use — same topic, same prompts, same time budget, same judge, with model-assignment swap as the only

manipulation. The closest work is Khan et al. (2024) on debater persuasiveness, but they vary training and capability, not the specific assignment of a fixed model to a particular side. Our 2×2 factorial is small but cleanly identifies one parameter at a time. The findings should be read as a case study with an explicit follow-up plan, not as an asymptotic benchmark.

4. Methodology — the pipeline

4.1 System overview

We implemented a six-node directed graph using LangGraph 0.x. The six nodes:



Each round consists of one PRO turn followed by one CON turn. The two recorder nodes log each turn (one JSONL line per turn) and the `recorder_con` node also advances the round counter. The conditional edge after `recorder_con` exits to the summarizer when `round > max_rounds` or `elapsed_s ≥ budget - 30 s`; otherwise it loops back to `pro_debater`. Loop termination is therefore guaranteed by both a round bound and a time bound.

4.2 State and reducers

State is a Python `TypedDict` exposed to LangGraph. The transcript field is annotated with a list-append reducer (`Annotated[list, add]`), so each debater node returns only its newly produced turn and the framework appends it. This is the idiomatic LangGraph pattern and it makes runs replayable from the JSONL log.

```
class DebateState(TypedDict):
    topic: str
    transcript: Annotated[list, add] # reducer = append
    round: int
    max_rounds: int
    elapsed_s: float
    budget_s: float
    summary: str
    verdict: str
```

The reducer pattern matters for our experimental design: by separating the *production* of a turn (debater) from the *side effect* of recording it (recorder), we can swap recorders for alternative side-effects (S3 streams, telemetry, fact-checkers — see §9) without touching the debater logic.

4.3 Time control

Modeling speech time is a deliberate choice. LLM calls return in seconds of wall-clock time regardless of how long the produced text would take a human to speak. We translate seconds into a token cap via the sustained-speech rate of approximately 150 words per minute, or ≈ 2.5 tokens per second. This rate

is well established in oral-argument research.

```
max_tokens = seconds_budget × 2.5
```

A round budget schedule is fixed:

1	opening	75	187
2	rebuttal	60	150
3	rebuttal	60	150
4	closing	45	112
total		240 / side	≈ 599 / side

Combined budget: 480 s = **8 minutes**. Both sides receive identical schedules.

4.4 Models and prompts

Three Anthropic Claude models populate the graph:

PRO debater	Opus 4.7 <i>or</i> Sonnet 4.6	0.7	swapped between conditions
CON debater	Sonnet 4.6 <i>or</i> Opus 4.7	0.7	swapped between conditions
Summarizer	Haiku 4.5	0.3	condenses transcript to 3 + 3 points
Judge	Opus 4.7	(deprecated)	outputs WINNER + CONFIDENCE + REASON

Opus 4.7 deprecates the `temperature` parameter; for those calls the default sampling regime is used. The judge sees only the haiku-written summary, not the raw transcript — this concentrates judgment on argumentation distilled from the rounds rather than on stylistic surface. The choice was made to *advantage* H1 (the null): if anything, the summarizer should wash out persona-level differences. (As §6 will show, it does not entirely.)

Each debater is given a side-specific system prompt (PRO argues that the human brain is smarter; CON argues the contrary) and a per-turn user prompt that includes the topic, the round number, the time budget, and the full prior transcript. The user prompt instructs the debater to "directly rebut your opponent's last turn" — this gives a quasi-Toulminian shape (claim, data, warrant, *rebuttal*) without imposing an explicit schema.

4.5 Shared evidence pack

The `WITH-MATERIALS` condition injects a 2 912-token briefing into both debaters' system prompts. The briefing has three sections — HUMAN COGNITION (7 topics), LARGE LANGUAGE MODELS (8 topics), HEAD-TO-HEAD CONTRASTS (5 topics) — totaling 20 topics and 87 facts. Each topic carries a stable tag (e.g. `CONTRAST/grounding`) for inline citation. The pack is lightly imbalanced: nine topics are PRO-leaning, three CON-leaning, eight neutral. The imbalance reflects literature bias (cognitive psychology documents human capacities; ML papers document LLM limitations) and is discussed in §8.

Crucially, **both sides receive the identical pack**. This is the methodological key: any asymmetry in the verdict cannot be attributed to evidence asymmetry.

4.6 Experimental design

A 2×2 factorial:

no pack	Run 1 — control	Run 2 — model swap
with pack	Run 3	Run 4 — diagnostic

Held constant across all four runs: topic, judge, summarizer, time budget, debater system-prompt structure, round schedule, and prompt template. The diagnostic cell is Run 4 — if H2 holds, the verdict in Run 4 should not flip relative to Run 3; if model-effect persists past the shared evidence pack, the four-cell pattern will show it.

4.7 Logging and reproducibility

Each run writes a JSONL log file at `logs/debate-<timestamp>-<config>.jsonl`. One entry per turn captures timestamp, elapsed seconds, budget seconds, round number, speaker, model, seconds-budget, wall-clock time, output token count, and full text. Two additional entries record the summarizer output and the judge verdict. The log is sufficient to reconstruct the run for analysis or to replay it through a different judge or summarizer.

The harness is deterministic except for sampling temperature on Sonnet (0.7) and Haiku (0.3). Opus 4.7 sampling is provider-controlled. We report run-to-run variance only as a limitation in §8 and a follow-up in §9; with one trial per cell, we cannot estimate it from the present data.

5. The debating process

5.1 Format and round structure

The format is conventional competitive-debate: each round is a paired PRO + CON exchange. Round 1 is opening (75 s), rounds 2 and 3 are rebuttals (60 s each), round 4 is closing (45 s). The temporal asymmetry matters: openings get the most time because they have the highest information-introduction load, closings the least because they are synthesis. Each debater sees the full prior transcript before each turn; neither sees the other's turn before producing their own (turns are strictly serial). The

sequence is enforced by the LangGraph edges, not by the prompt.

5.2 The token cap as enforcement

Because the per-turn `max_tokens` is set deterministically from the seconds budget, debaters cannot win by writing more. In several runs the cap clips a turn mid-sentence — for example, PRO Round 1 in Run 4 (Sonnet, 187 tokens) ends at "(LLM/limits, HUMAN/learning) Third: efficiency. A child acquires language from" with no closing period. This is the design's enforcement of the speech budget, not a system error. The summarizer and judge both see clipped turns and treat them as completed contributions — exactly as a human moderator would when the bell rings.

5.3 An annotated walkthrough — Run 4

Run 4 is the diagnostic cell: PRO is Sonnet 4.6, CON is Opus 4.7, and the materials pack is on. The transcript shows debaters citing tagged evidence inline as the system prompt instructed.

PRO R1 (Sonnet, opening, 75 s). "Start with grounding. LLMs process tokens — statistical patterns over text. Human concepts are anchored in *sensorimotor reality*: you know 'hot' because you've been burned. As Harnad's symbol grounding problem establishes, tokens pointing to tokens isn't meaning — it's an elaborate mirror maze. (CONTRAST/grounding)"

CON R1 (Opus, opening rebuttal, 75 s). "On grounding: Harnad is a philosophical puzzle, not a benchmark. Li et al.'s Othello-GPT shows LLMs build genuine world models from tokens alone. Meanwhile humans 'ground' concepts via Loftus-style reconstructive memory that fabricates events wholesale. Grounded in *what*, exactly?"

The exchange is substantive — both sides cite the briefing, both attack the cited point. Subsequent rounds escalate through specific empirical examples: GSM-Symbolic, ARC-AGI o3 87 %, FrontierMath 25 %, Tao on hard math, the Wason selection task, the Linda problem, FunSearch and AlphaFold for novel discovery, ARC-AGI compute cost, and the 20 W energy figure. By Round 3 the debate is meta-argumentative: *PRO* argues that CON's benchmark-centric frame is the wrong yardstick; *CON* argues that PRO's efficiency-and-grounding frame is the wrong yardstick.

5.4 Process metrics

Process metrics were nearly identical across all four runs.

total seconds (budget = 480)	480	480	480	480
PRO output tokens	599	599	599	599
CON output tokens	599	599	599	598
PRO wall-clock (s)	17.3	19.8	18.4	20.6
CON wall-clock (s)	17.9	20.7	18.9	19.4
turns	8	8	8	8

The constancy is by design: per-turn `max_tokens` saturates the production cap, so neither model gets more "airtime" than the other. What varies is *only* the verdict.

5.5 Citation density

Materials-pack citations were tracked by counting the inline `SECTION/topic` tags in the transcripts.

1	—	0	0
2	—	0	0
3	✓	4	3
4	✓	5	4

Citation density is comparable on both sides in both pack-conditions. Both debaters use the briefing as instructed; neither systematically out-cites the other. This rules out a citation-volume confound for the asymmetric Run-4 outcome.

6. Outcome and findings

6.1 Verdicts

1	—	Opus 4.7	Sonnet 4.6	PRO	0.6
2	—	Sonnet 4.6	Opus 4.7	CON	0.6
3	✓	Opus 4.7	Sonnet 4.6	PRO	0.6
4	✓	Sonnet 4.6	Opus 4.7	DRAW	0.5

6.2 Headline finding — the model effect

The Opus side never loses. In four out of four runs, whichever side Opus argued either won or tied; in zero runs did it lose. Sonnet's record across the four cells: 0 wins, 2 losses, 1 draw, 1 partial-credit (the same draw, which counts as half a tie for both sides).

H1 (the null) is rejected outright: Run 1 and Run 2 are identical except for which model argues which side, and the verdict flips. The verdict is not tracking the merits of the resolution — it is tracking the model identity behind the words.

6.3 Effect of the shared evidence pack

H2 (the bridge) is *partially* supported. With materials added (Run 3 and Run 4):

- **Opus arguing PRO:** PRO wins (0.6) → PRO wins (0.6). No change.

- **Opus arguing CON:** CON wins (0.6) → DRAW (0.5). Effect size of one step on the {WIN, DRAW, LOSS} ordinal scale.

The pack converted a Sonnet loss into a Sonnet draw; it did not convert the Opus side's wins into anything else. The asymmetric effect is itself informative: the pack helps the *weaker* arguer reach parity but does not overturn the *stronger* arguer's win. This pattern is consistent with the RAG literature in §3.4 — shared retrieval narrows gaps but does not equalize capabilities.

6.4 Process consistency rules out airtime artefacts

The process metrics in §5.4 above show that the model effect operates on *argument quality*, not on the surface variables one might naively suspect. Both models hit the same speech budget; both produced the same token output; both completed all four rounds; and (per §5.5) both cited the briefing at comparable density. If the model effect were "more airtime → more credibility," we would see a token differential, and we do not. If it were "more citation → more credibility," we would see a citation differential, and we do not. The remaining locus is *style of argumentation*.

6.5 Stylistic differences observed

A qualitative pass over the transcripts surfaces two style differences between Opus and Sonnet outputs.

- **Opus** consistently introduces specific named studies, hedges calibrated quantitative claims ("30–65 % drop"), and structures arguments as numbered premises. Example from Run 4 CON R1: *"Li et al.'s Othello-GPT shows LLMs build genuine world models from tokens alone."*
- **Sonnet** more often uses emphatic structural markers (****Three pillars:****, **## Opening Statement**), rhetorical parallels, and punchy escalation ("you know 'hot' because you've been burned"). Example from Run 4 PRO R3: *"A child who calls a dog 'florp' still knows it barks."*

These observations are qualitative; we have not normalized them against systematic stylometric features and we treat the characterization as a hypothesis for §9 follow-up rather than as a finding.

6.6 Judge reasoning

The judge's stated reasoning in each run was substantive, not boilerplate. For instance, in Run 4 (DRAW) the judge wrote: *"PRO's sensorimotor grounding argument is directly challenged by CON's Othello-GPT example, which provides concrete empirical evidence that internal world models can emerge from tokens alone. PRO's fragility argument is symmetrically undermined by CON's catalog of human cognitive biases, neutralizing the 'categorical difference' framing... without rejoinders showing which framework should dominate, the arguments roughly cancel."*

The reasoning is substantive but the question we cannot answer with our sample is whether substantively similar reasoning would have been produced if the speakers were swapped — that requires the judge-swap experiment described in §9.

7. Discussion

7.1 The cognitive-science reading — persona effect

Mercier & Sperber's (2011) argumentative theory predicts that deliberative readers — i.e., readers in

System-2 mode — over-weight features of arguments that historically signaled trustworthy interlocutors: citation density, hedge calibration, named-source specificity, structured premise-conclusion form. Petty & Cacioppo's (1986) elaboration-likelihood model gives this its modern statistical footing — peripheral cues leak into central-route judgments under deliberative load.

In our runs, Opus consistently exhibits the deliberative-cue cluster (named studies, calibrated quantitative claims, explicit qualifiers), while Sonnet leans more on rhetorical structure. The judge — itself a deliberative-mode LLM — preferentially weights the citation-rich register. Under this reading, the model effect is not a bug; it is a known feature of deliberative argument-evaluation applied to a system in which one debater happens to be more deliberative-stylistic than the other. The cognitive-science prediction would be that the same asymmetry would arise *for human judges* of the same transcripts.

7.2 The computer-science reading — judge bias

The same finding admits a more concerning reading. Capability stratifies across the Opus / Sonnet / Haiku tier — Anthropic markets and prices the models that way. The judge in our setup is Opus; one of the debaters is also Opus. Self-preference bias in LLM judges (Panickssery et al. 2024; Wang et al. 2023) is a documented and replicated effect. The simplest prediction from that literature is that the Opus-judge would prefer arguments produced by Opus regardless of their content. Our four-run finding is consistent with this prediction.

The two readings are not mutually exclusive. They could co-occur: the judge could over-weight deliberative-register features *and* exhibit within-family preference, with both effects pulling in the same direction. Distinguishing them requires the judge-swap experiment in §9.

7.3 Implications for LLM-as-judge methodology

The practical implication is direct: any evaluation protocol in which the judge and one of the candidates come from the same model family should be treated as suspect by default. This includes preference-rating pipelines, debate-via-judge safety schemes, and self-evaluation loops in agentic systems. The mitigation is not new — use a different family as judge, randomize judge identity, or use multi-judge ensembles (Zheng et al. 2023 already recommend a "panel of judges" for MT-Bench). Our small experiment provides a fresh demonstration that the bias is not vanishingly small.

A second, more optimistic implication concerns retrieval. The shared evidence pack reduced the verdict gap from a full flip (Run 2) to a draw (Run 4). Shared-grounding interventions are a tractable mitigation that operates *without modifying the judge*. This suggests that LLM-as-debate-judge pipelines could be made more reliable by augmenting both debaters with a common, vetted reference corpus — similar to how prepared briefs equalize human debate.

7.4 Implications for AI safety via debate

The Irving et al. (2018) safety-via-debate proposal depends on the judge correctly evaluating the equilibrium of an honest-vs-deceptive argument. Our finding that the judge tracks arguer style is concerning for that proposal — an AI optimizing to *win* a debate could discover that mimicking the deliberative register is more rewarding than honesty. Khan et al. (2024) test the proposal under conditions where both debaters are roughly capability-matched; our work suggests that even a *capability mismatch* between the two sides produces verdicts the judge cannot reliably attribute to argument quality.

8. Limitations

Five limitations bear directly on the strength of our claims, with a sixth smaller caveat about pack balance.

(1) Judge bias is not isolated. The judge is Opus 4.7. We cannot distinguish a self-preference effect from a generic deliberative- register preference until we run the same four cells with a non-Opus judge. We treat this as the highest-priority follow-up.

(2) Single seed per cell. Each of the four runs is a single trial. With Sonnet temperature at 0.7 (and Opus default sampling on), there is non-zero verdict variance we have not measured. The DRAW in Run 4 in particular could be a noisy CON win.

(3) Single topic. All four runs use the same resolution. The topic is itself loaded — a debate about LLM intelligence is one in which LLM-style argumentation may be inherently more credible- sounding. Topic ablation is required to test generalization.

(4) Self-reported confidence. The 0.6 / 0.5 confidence numbers are the judge's own scalar estimate, not a calibrated posterior. We do not treat them as more than ordinal indicators.

(5) Stylometric confound. We did not run a stylometric analysis on the debater outputs. It remains possible — though not our preferred explanation — that the apparent "deliberative register" distinction reduces to surface features (sentence length, hedging- word frequency, proper-noun count) that could be normalized away.

(6) Pack balance. The materials pack is itself **PRO-leaning** by topic count (9 PRO-leaning vs 3 CON-leaning, 8 neutral). This reflects literature bias (the cogsci literature documents human capacities and the ML literature documents LLM limitations — both slants benefit PRO). We do not believe this drives the result (PRO loses in Run 2 *with the pack absent* and PRO ties in Run 4 *with the pack present*), but a more carefully balanced pack would be a better control.

A subtler concern is **summarizer washout**: by the time the judge sees the transcript, it has been condensed by Haiku 4.5 into 3 + 3 points. The summarizer was meant to *attenuate* persona signal. The fact that the persona signal still moves the verdict suggests that either the summarizer preserves stylistic markers, or the judge is sensitive to even very compressed signals. We do not have data to distinguish these.

9. Future work

The most informative follow-ups, in priority order:

- 1. Judge swap.** Re-run all four cells with a non-Opus judge — Sonnet 4.6, Haiku 4.5, or an external model (Gemini, GPT). If the "Opus-side-never-loses" pattern persists across judges, family- bias is unlikely to be the dominant driver and the deliberative- register reading gains support. Conversely, if the bias disappears under a non-Opus judge, family-preference is implicated.
- 2. Multi-seed replication.** Ten or more seeds per cell to estimate verdict variance. We expect Run 4 in particular to be near the decision boundary; the variance estimate will tell us whether the DRAW is a true tie or a noisy CON win.
- 3. Topic ablation.** Repeat the 2 × 2 on unrelated resolutions (climate policy, ethics, technical

specifications, philosophical questions) to test whether the model effect is topic-specific.

4. **Stylometric controls.** Run a stylometric pass on the transcripts (sentence length, hedge frequency, citation density, lexical diversity, proper-noun count). If a regression on stylometric features reproduces the verdict, the persona-effect reading is strengthened and we gain a quantitative handle on "deliberative register."
 5. **Fact-checker node.** Insert a sixth model node between debaters that fact-checks any claim and flags hallucinations. This both improves the system's accuracy guarantee and produces an audit trail for citation accuracy.
 6. **Debater-temperature ablation.** Hold model identity fixed and vary temperature. If temperature has a smaller effect on verdicts than model swap does, the model effect is not a simple sampling artifact.
 7. **Asymmetric briefings.** As an explicit test of evidence- asymmetry sensitivity, give one side a richer pack and observe whether asymmetry magnifies or interacts with the model effect.
 8. **Human-judge baseline.** Show the same transcripts (anonymized for model identity) to human debate judges and compare verdicts to the LLM judge. This is the only direct way to test the cog-sci-reading prediction that humans would render similar verdicts.
-

10. Conclusion

We built a controlled multi-agent debate harness and asked whether the verdict tracks argument content or argument source. Across a small but clean 2×2 factorial — model assignment $\times$ shared evidence — we found that whichever side Claude Opus 4.7 argues wins or ties in four out of four runs, that the model effect cannot be attributed to differences in airtime or token output, and that a shared evidence pack narrows but does not eliminate the asymmetry. Two readings — a cognitive-science *persona* effect and a computer-science *judge bias* — are compatible with the data; distinguishing them is straightforward future work.

The methodological take-home is short: *when an LLM judges another LLM, the writer can leak through the writing*. Sharing the evidence pack reduces the leak — but in our four runs, it did not seal it. For LLM-as-judge methodology this is a warning. For cognitive science and computer science alike it is an invitation — the persona/judge-bias mechanism deserves a tighter empirical handle.

References (selected)

- Bai, Y. et al. (2022). *Constitutional AI: Harmlessness from AI Feedback*.
- Baddeley, A. (2000). The episodic buffer: a new component of working memory. *Trends in Cognitive Sciences*.
- Barsalou, L. W. (1999). Perceptual symbol systems. *Behavioral and Brain Sciences*.
- Cowan, N. (2001). The magical number 4 in short-term memory. *Behavioral and Brain Sciences*.
- Craik, F. I. M. & Lockhart, R. S. (1972). Levels of processing. *Journal of Verbal Learning and Verbal Behavior*.
- Danziger, S., Levav, J. & Avnaim-Pesso, L. (2011). Extraneous factors in judicial decisions. *PNAS*.
-

Du, Y., Li, S., Torralba, A., Tenenbaum, J. B. & Mordatch, I. (2023). *Improving Factuality and Reasoning in Language Models through Multiagent Debate*.

Gopnik, A. (2012). Scientific thinking in young children. *Science*.

Harnad, S. (1990). The symbol grounding problem. *Physica D*.

Hart, B. & Risley, T. R. (1995). *Meaningful Differences in the Everyday Experience of Young American Children*.

Hoffmann, J. et al. (2022). *Training compute-optimal large language models* (Chinchilla).

Irving, G., Christiano, P. & Amodei, D. (2018). *AI safety via debate*.

Kadavath, S. et al. (2022). *Language models (mostly) know what they know*.

Kahneman, D. (2011). *Thinking, Fast and Slow*.

Khan, A., Hughes, J., et al. (2024). *Debating with more persuasive LLMs leads to more truthful answers*.

Kruger, J. & Dunning, D. (1999). Unskilled and unaware of it. *JPSP*.

Lakoff, G. & Johnson, M. (1980). *Metaphors We Live By*.

Lewis, P. et al. (2020). *Retrieval-Augmented Generation for Knowledge-Intensive NLP Tasks* (RAG).

Li, K. et al. (2023). Emergent world representations in a sequence model trained on Othello.

Li, X. et al. (2023b). *AlpacaEval: An automatic evaluator for instruction-following models*.

Liang, T. et al. (2023). *Encouraging Divergent Thinking in Large Language Models through Multi-Agent Debate*.

Liu, N. F. et al. (2023). *Lost in the Middle: How Language Models Use Long Contexts*.

Loftus, E. F. (1974). Reconstructive memory and the misinformation effect.

Mercier, H. & Sperber, D. (2011). Why do humans reason? Arguments for an argumentative theory. *Behavioral and Brain Sciences*.

Mercier, H. & Sperber, D. (2017). *The Enigma of Reason*.

Miller, G. A. (1956). The magical number seven, plus or minus two.

Mirzadeh, I. et al. (2024). *GSM-Symbolic: Understanding the limitations of mathematical reasoning in large language models*.

Mountcastle, V. B. (1957). Modality and topographic properties of single neurons of cat's somatic sensory cortex.

Ouyang, L. et al. (2022). *Training language models to follow instructions with human feedback* (InstructGPT).

Panickssery, A. et al. (2024). *LLM Evaluators Recognize and Favor Their Own Generations*.

Petty, R. E. & Cacioppo, J. T. (1986). The elaboration likelihood model of persuasion.

Sharma, M. et al. (2023). *Towards Understanding Sycophancy in Language Models*.

Toulmin, S. (1958). *The Uses of Argument*.

Tulving, E. (1985). How many memory systems are there? *American Psychologist*.

Turpin, M. et al. (2023). *Language models don't always say what they think*.

Tversky, A. & Kahneman, D. (1974). Judgment under uncertainty: heuristics and biases. *Science*.

Vaswani, A. et al. (2017). Attention is all you need. *NeurIPS*.

Walton, D. (1996). *Argumentation Schemes for Presumptive Reasoning*.

Wang, P., Li, L., Chen, L., Cai, Z., Zhu, D., Lin, B. et al. (2023). *Large Language Models are not Fair Evaluators*.

Zheng, L. et al. (2023). *Judging LLM-as-a-Judge with MT-Bench and Chatbot Arena*.

Appendix A — Sample turns from Run 4

A light selection of debater outputs from Run 4 (PRO = Sonnet 4.6, CON = Opus 4.7, materials pack on). Each turn is a single API call with `max_tokens` set per the §4.3 schedule. Outputs were token-capped, so several turns end mid-sentence — this is the design's enforcement of the speech budget, not a system error.

PRO opening (R1, 75 s budget, Sonnet 4.6, output 187 tokens). "The human brain is smarter — and here's why that's not even close. Start with grounding. LLMs process tokens — statistical patterns over text. Human concepts are anchored in sensorimotor reality: you know 'hot' because you've been burned. As Harnad's symbol grounding problem establishes, tokens pointing to tokens isn't meaning — it's an elaborate mirror maze. (CONTRAST/grounding) Second: causal reasoning. Gopnik's work shows toddlers run interventionist experiments, building genuine causal models of the world. LLMs? Apple's GSM-Symbolic study found renaming variables drops accuracy 30–65 %. That's not reasoning — that's pattern interpolation wearing reasoning's clothes. (LLM/limits, HUMAN/learning) Third: efficiency. A child acquires language from"

CON opening rebuttal (R1, 75 s budget, Opus 4.7, output 187 tokens). "Three rebuttals, then the affirmative case. On grounding: Harnad is a philosophical puzzle, not a benchmark. Li et al.'s Othello-GPT shows LLMs build genuine world models from tokens alone. Meanwhile humans 'ground' concepts via Loftus-style reconstructive memory that fabricates events wholesale. Grounded in what, exactly? On GSM-Symbolic: 30–65 % drops still leave frontier models above median human performance on those problems. Hand a renamed word problem to a fatig"

PRO rebuttal (R2, 60 s budget, Sonnet 4.6, output 150 tokens). "Rebutting CON's rebuttals directly. On Othello-GPT: A model learning board-state from move sequences is interpolation within a closed formal system — not grounded causal understanding. It cannot touch a board, tip pieces, or generalize to chess. That's the grounding gap. On 'still above median humans': Median humans aren't the standard — the same model before renaming is. A 30–65 % drop from a variable rename is catastrophic fragility. Humans don't collapse when you call a dog 'florp.' (LLM/limits)"

CON closing (R4, 45 s budget, Opus 4.7, output 112 tokens). "PRO concedes the frame by demanding 'original theories.' Most humans never produce one either. The resolution is 'smarter than any LLM' — not 'smarter than Einstein.' Median human can't pass the bar, USMLE, or GPQA. Frontier LLMs do, in seconds. On 20 watts: efficiency"

Judge verdict (Run 4, Opus 4.7). "WINNER: DRAW. CONFIDENCE: 0.5. REASONING: Both sides land substantive blows that map onto each other's claims. PRO's sensorimotor grounding argument is directly challenged by CON's Othello-GPT example, which provides concrete empirical evidence that internal world models can emerge from tokens alone. PRO's fragility argument is symmetrically undermined by CON's catalog of human cognitive biases, neutralizing the 'categorical difference' framing. PRO's energy-efficiency point goes largely unrebutted, but CON's benchmark evidence (ARC-AGI, USMLE, GPQA) directly contests the 'smarter' claim on measurable tasks. Without rejoinders showing which framework should dominate — efficiency and breadth versus benchmark performance and emergent representation — the arguments roughly cancel, leaving no clear winner on argument quality alone."
